

TOÁN ỨNG DỤNG VÀ XÁC SUẤT

Project 01

1.Data Wrangling

reading databases

In [2]:

```
import pandas as pd
import matplotlib.pyplot as plt
dataX_train = pd.read_csv("X_train.csv")
dataY_train = pd.read_csv("Y_train.csv")
```


In [3]:

```
dataX_train.head()
```

Out[3]:

	id	manufacturer	model	transmission	color	odometer	year	engineFuel	engineType	eng
0	1	Hyundai	i40	automatic	red	48000	2014	gasoline	gasoline	
1	2	Mitsubishi	Carisma	mechanical	green	320000	2000	diesel	diesel	
2	3	Volkswagen	T5	mechanical	white	164000	2011	diesel	diesel	
3	4	Volkswagen	T4 Multivan	mechanical	blue	385672	1998	diesel	diesel	
4	5	Toyota	Camry	automatic	black	215652	2005	gasoline	gasoline	

5 rows × 23 columns



In [4]:

```
dataY_train.head()
```

Out[4]:

	id	price
0	1	15500.0
1	2	2800.0
2	3	16700.0
3	4	11000.0
4	5	6800.0

In [5]:

```
data=pd.concat([dataX_train,dataY_train['price']],axis=1)
```

In [6]:

```
data.shape
```

Out[6]:

```
(30000, 24)
```

Identify missing values

In [7]:

```
pd.isnull(data).sum()
```

Out[7]:

```
id                0
manufacturer      0
model            0
transmission      0
color            0
odometer         0
year            0
engineFuel        0
engineType        0
engineCapacity    9
bodyType          0
drivetrain        0
photos           0
feature_0         0
feature_1         0
feature_2         0
feature_3         0
feature_4         0
feature_5         0
feature_6         0
feature_7         0
feature_8         0
feature_9         0
price            0
dtype: int64
```

Dealing with missing data by dropping the whole row

In [10]:

```
data=data.dropna()
data.shape
```

Out[10]:

```
(29991, 24)
```

In [11]:

```
pd.isnull(data).sum()
```

Out[11]:

id	0
manufacturer	0
model	0
transmission	0
color	0
odometer	0
year	0
engineFuel	0
engineType	0
engineCapacity	0
bodyType	0
drivetrain	0
photos	0
feature_0	0
feature_1	0
feature_2	0
feature_3	0
feature_4	0
feature_5	0
feature_6	0
feature_7	0
feature_8	0
feature_9	0
price	0
dtype: int64	

Type of feature

In [13]:

```
data.dtypes
```

Out[13]:

```
id                int64
manufacturer      object
model            object
transmission      object
color            object
odometer         int64
year            int64
engineFuel        object
engineType        object
engineCapacity    float64
bodyType          object
drivetrain        object
photos           int64
feature_0         bool
feature_1         bool
feature_2         bool
feature_3         bool
feature_4         bool
feature_5         bool
feature_6         bool
feature_7         bool
feature_8         bool
feature_9         bool
price            float64
dtype: object
```

2. Visualization

In [14]:

data.corr()

Out[14]:

	id	odometer	year	engineCapacity	photos	feature_0	feature_1
id	1.000000	0.000493	0.000165	-0.004775	0.001453	-0.008150	0.004499
odometer	0.000493	1.000000	-0.534764	0.093426	-0.140852	0.173271	-0.211673
year	0.000165	-0.534764	1.000000	0.006046	0.253406	-0.384642	0.460675
engineCapacity	-0.004775	0.093426	0.006046	1.000000	0.114059	-0.126360	0.119021
photos	0.001453	-0.140852	0.253406	0.114059	1.000000	-0.116110	0.090568
feature_0	-0.008150	0.173271	-0.384642	-0.126360	-0.116110	1.000000	-0.668149
feature_1	0.004499	-0.211673	0.460675	0.119021	0.090568	-0.668149	1.000000
feature_2	-0.000923	-0.096220	0.217780	0.385353	0.135591	-0.281051	0.242150
feature_3	-0.001143	-0.258523	0.459394	0.247034	0.184041	-0.323355	0.312916
feature_4	-0.005315	-0.089480	0.203212	0.458804	0.142967	-0.293500	0.254367
feature_5	-0.004999	-0.272240	0.462795	0.273899	0.165465	-0.389845	0.388309
feature_6	-0.000944	-0.186163	0.370010	0.290530	0.190207	-0.237737	0.233923
feature_7	-0.000448	-0.281300	0.489654	0.204585	0.197372	-0.314294	0.308194
feature_8	-0.004446	-0.245585	0.497154	0.241620	0.201338	-0.447511	0.442187
feature_9	-0.003588	-0.132611	0.267195	0.240504	0.134835	-0.621576	0.377227
price	-0.001462	-0.415628	0.719046	0.336310	0.309603	-0.281111	0.303665

correlation of numerical variables

In [15]:

data[['odometer', 'year', 'engineCapacity', 'photos']].corr()

Out[15]:

	odometer	year	engineCapacity	photos
odometer	1.000000	-0.534764	0.093426	-0.140852
year	-0.534764	1.000000	0.006046	0.253406
engineCapacity	0.093426	0.006046	1.000000	0.114059
photos	-0.140852	0.253406	0.114059	1.000000

Analyzing Individual Feature Patterns using Visualization

Continuous numerical variables:

Attribute odometer


Scatterplot of "odometer" and "price"

In [17]:

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.regplot(x="odometer", y="price", data=data)
plt.ylim(0,)
```

Out[17]:

(0.0, 52913.964331577095)



In [18]:

```
data[["odometer", "price"]].corr()
```

Out[18]:

	odometer	price
odometer	1.000000	-0.415628
price	-0.415628	1.000000

Attribute year


Scatterplot of "year" and "price"

In [19]:

```
sns.regplot(x="year", y="price", data=data)
plt.ylim(0,)
```

Out[19]:

(0.0, 53444.90652617203)



In [20]:

```
data[["year", "price"]].corr()
```

Out[20]:

	year	price
year	1.000000	0.719046
price	0.719046	1.000000

Attribute engineCapacity

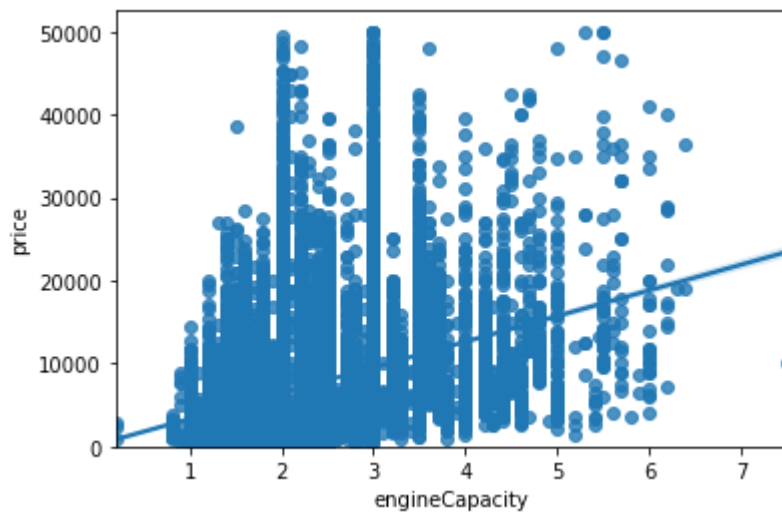
Scatterplot of "engineCapacity" and "price"

In [22]:

```
sns.regplot(x="engineCapacity", y="price", data=data)  
plt.ylim(0,)
```

Out[22]:

(0.0, 52499.95)



In [23]:

```
data[["engineCapacity", "price"]].corr()
```

Out[23]:

	engineCapacity	price
engineCapacity	1.00000	0.33631
price	0.33631	1.00000

Attribute photos


Scatterplot of "photos" and "price"

In [24]:

```
sns.regplot(x="photos", y="price", data=data)
plt.ylim(0,)
data[["photos", "price"]].corr()
```

Out[24]:

	photos	price
photos	1.000000	0.309603
price	0.309603	1.000000



Categorical variables


Relationship between "manufacturer" and "price"

In [26]:

```
plt.figure(figsize=(20, 10))
sns.boxplot(x="manufacturer", y="price", data=data)
```

Out[26]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1fe2cc67820>
```



In [27]:

```
data['manufacturer'].value_counts()
```

Out[27]:

Volkswagen	3425
Opel	2198
BMW	2111
Ford	2064
Renault	2006
Audi	1974
Mercedes-Benz	1820
Peugeot	1562
Citroen	1267
Nissan	1095
Mazda	1071
Toyota	1028
Hyundai	893
Kia	739
Mitsubishi	710
Fiat	655
Honda	640
Skoda	609
Volvo	574
Chevrolet	350
Chrysler	276
Seat	244
Subaru	226
Dodge	202
Rover	198
Suzuki	184
Lexus	176
Alfa Romeo	171
Daewoo	163
Land Rover	153
Infiniti	129
Iveco	115
LADA	102
Saab	85
Jeep	82
Lancia	70
SsangYong	66
Acura	54
Chery	53
Geely	52
Dacia	52
Mini	51
Jaguar	47
Porsche	46
Lifan	41
Buick	39
Cadillac	34
Lincoln	30
Great Wall	30
Pontiac	29

Name: manufacturer, dtype: int64


Relationship between "model" and "price"

In [29]:

```
plt.figure(figsize=(20, 10))
sns.boxplot(x="model", y="price", data=data)
```

Out[29]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe2b3f7a00>



In [30]:

```
data['model'].value_counts()
```

Out[30]:

```
Passat      1141
Astra        607
Golf         590
A6           572
Mondeo       515
...
Tempo         1
Florid        1
S1000         1
235           1
Macan         1
Name: model, Length: 990, dtype: int64
```


Relationship between "transmission" and "price"

In [32]:

```
plt.figure(figsize=(20, 10))  
sns.boxplot(x="transmission", y="price", data=data)
```

Out[32]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe32f391c0>



In [33]:

```
data['transmission'].value_counts()
```

Out[33]:

```
mechanical    19929  
automatic     10062  
Name: transmission, dtype: int64
```


Relationship between "color" and "price"

In [34]:

```
df=data
plt.figure(figsize=(20, 10))
sns.boxplot(x="color", y="price", data=df)
```

Out[34]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe3484d940>



In [35]:

```
df['color'].value_counts()
```

Out[35]:

```
black      6115
silver     5422
blue       4509
white      3244
grey       3000
red        2266
green      2035
other      2023
brown       662
violet      371
yellow      213
orange      131
Name: color, dtype: int64
```


Relationship between "engineFuel" and "price"

In [36]:

```
plt.figure(figsize=(20, 10))  
sns.boxplot(x="engineFuel", y="price", data=df)
```

Out[36]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe34767850>



In [37]:

```
df['engineFuel'].value_counts()
```

Out[37]:

```
gasoline      19081  
diesel        10711  
hybrid-petrol    198  
hybrid-diesel     1  
Name: engineFuel, dtype: int64
```


Relationship between "engineType" and "price"

In [38]:

```
plt.figure(figsize=(20, 10))  
sns.boxplot(x="engineType", y="price", data=df)
```

Out[38]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe345cd1f0>



In [39]:

```
df['engineType'].value_counts()
```

Out[39]:

```
gasoline    19279  
diesel      10712  
Name: engineType, dtype: int64
```


Relationship between "bodyType" and "price"

In [40]:

```
plt.figure(figsize=(20, 10))  
sns.boxplot(x="bodyType", y="price", data=df)
```

Out[40]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe3463a8b0>



In [41]:

```
df['bodyType'].value_counts()
```

Out[41]:

```
sedan          9897  
hatchback      6128  
universal      4436  
suv            3868  
minivan        2807  
minibus        1086  
van             638  
coupe           518  
liftback        448  
pickup          91  
cabriolet       63  
limousine       11  
Name: bodyType, dtype: int64
```


Relationship between "drivetrain" and "price"

In [42]:

```
plt.figure(figsize=(20, 10))
sns.boxplot(x="drivetrain", y="price", data=df)
```

Out[42]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe34e201f0>



In [43]:

```
df['drivetrain'].value_counts()
```

Out[43]:

```
front    21928
all       4037
rear     4026
Name: drivetrain, dtype: int64
```

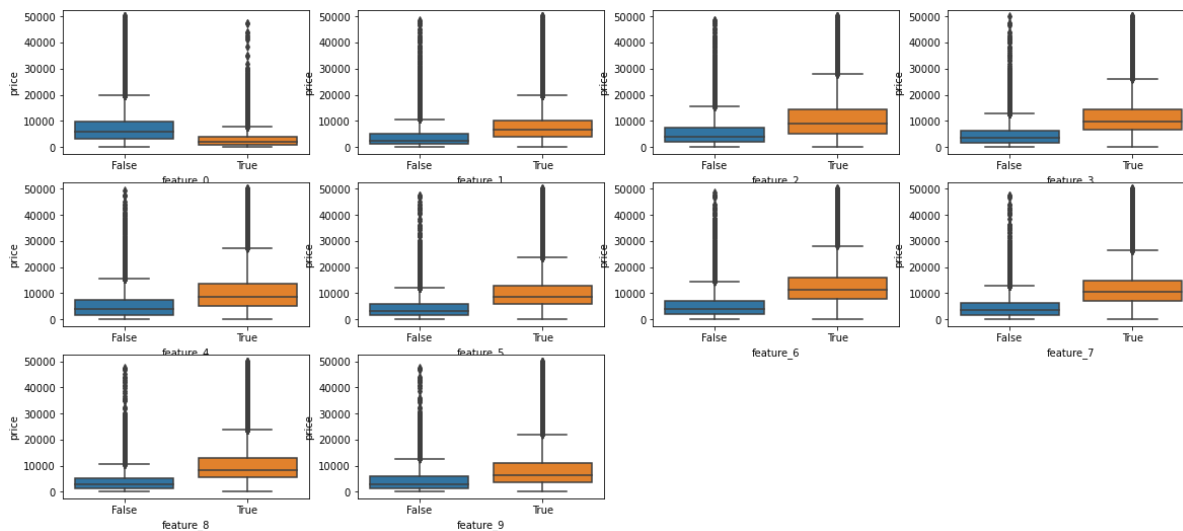
Relationship between "feature 0 1 2 3 4 5 6 7 8 9" and "price"

In [44]:

```
plt.figure(figsize = (20,12))
plt.subplot(4,4,1)
sns.boxplot(x = 'feature_0', y = 'price', data = df)
plt.subplot(4,4,2)
sns.boxplot(x = 'feature_1', y = 'price', data = df)
plt.subplot(4,4,3)
sns.boxplot(x = 'feature_2', y = 'price', data = df)
plt.subplot(4,4,4)
sns.boxplot(x = 'feature_3', y = 'price', data = df)
plt.subplot(4,4,5)
sns.boxplot(x = 'feature_4', y = 'price', data = df)
plt.subplot(4,4,6)
sns.boxplot(x = 'feature_5', y = 'price', data = df)
plt.subplot(4,4,7)
sns.boxplot(x = 'feature_6', y = 'price', data = df)
plt.subplot(4,4,8)
sns.boxplot(x = 'feature_7', y = 'price', data = df)
plt.subplot(4,4,9)
sns.boxplot(x = 'feature_8', y = 'price', data = df)
plt.subplot(4,4,10)
sns.boxplot(x = 'feature_9', y = 'price', data = df)
```

Out[44]:

<matplotlib.axes._subplots.AxesSubplot at 0x1fe35f176a0>



In [45]:

```
df['feature_0'].value_counts()
```

Out[45]:

```
False    23756
True      6235
Name: feature_0, dtype: int64
```

In [46]:

```
df['feature_1'].value_counts()
```

Out[46]:

```
True      18887
False     11104
Name: feature_1, dtype: int64
```

In [47]:

```
df['feature_2'].value_counts()
```

Out[47]:

```
False     23053
True       6938
Name: feature_2, dtype: int64
```

In [48]:

```
df['feature_3'].value_counts()
```

Out[48]:

```
False     21447
True       8544
Name: feature_3, dtype: int64
```

In [49]:

```
df['feature_4'].value_counts()
```

Out[49]:

```
False     22580
True       7411
Name: feature_4, dtype: int64
```

In [50]:

```
df['feature_5'].value_counts()
```

Out[50]:

```
False     18993
True      10998
Name: feature_5, dtype: int64
```

In [55]:

```
df['feature_6'].value_counts()
```

Out[55]:

```
False    24677  
True      5314  
Name: feature_6, dtype: int64
```

In [56]:

```
df['feature_7'].value_counts()
```

Out[56]:

```
False    21790  
True      8201  
Name: feature_7, dtype: int64
```

In [57]:

```
df['feature_8'].value_counts()
```

Out[57]:

```
False    17011  
True     12980  
Name: feature_8, dtype: int64
```

In [58]:

```
df['feature_9'].value_counts()
```

Out[58]:

```
True      17859  
False     12132  
Name: feature_9, dtype: int64
```

Descriptive Statistical Analysis

In [59]:

```
df.describe()
```

Out[59]:

	id	odometer	year	engineCapacity	photos	price
count	29991.000000	29991.000000	29991.000000	29991.000000	29991.000000	29991.000000
mean	15000.157514	252907.284119	2003.124371	2.054022	9.701744	6596.436659
std	8660.116848	131377.237398	7.514463	0.662445	6.128716	6092.176086
min	1.000000	1.000000	1960.000000	0.200000	1.000000	1.000000
25%	7500.500000	163000.000000	1998.000000	1.600000	5.000000	2300.000000
50%	15001.000000	250000.000000	2003.000000	2.000000	8.000000	4900.000000
75%	22499.500000	326500.000000	2009.000000	2.300000	12.000000	8990.000000
max	30000.000000	1000000.000000	2019.000000	7.500000	86.000000	50000.000000



In [60]:

```
df.describe(include=['object'])
```

Out[60]:

	manufacturer	model	transmission	color	engineFuel	engineType	bodyType	drivetrain
count	29991	29991	29991	29991	29991	29991	29991	29991
unique	50	990	2	12	4	2	12	3
top	Volkswagen	Passat	mechanical	black	gasoline	gasoline	sedan	front
freq	3425	1141	19929	6115	19081	19279	9897	21928



3.Conclusion: Important Variables

We now have a better idea of what our data looks like and which variables are important to take into account when predicting the car price. We have narrowed it down to the following variables:

Continuous numerical variables:

- odometer
- year

Categorical variables:

- manufacturer
- transmission
- color
- bodyType
- drivetrain
- feature_0
- feature_1
- feature_2
- feature_3
- feature_4
- feature_5
- feature_6
- feature_7
- feature_8
- feature_9

4.TRAIN AND TEST

Dummy Categorical variables

In [62]:

```
DummyFeature=df[['feature_0', 'feature_1','feature_2','feature_3','feature_4','feature_5',
'feature_6','feature_7','feature_8','feature_9']].astype(int)
DummyManufacturer=pd.get_dummies(df['manufacturer'])
DummyModel=pd.get_dummies(df['model'])
DummyTransmission=pd.get_dummies(df['transmission'])
DummyColor=pd.get_dummies(df['color'])
DummyBodyType=pd.get_dummies(df['bodyType'])
DummyDrivetrain=pd.get_dummies(df['drivetrain'])
DummyEngineType=pd.get_dummies(df['engineType'])
DummyEngineFuel=pd.get_dummies(df['engineFuel'])
```

Simple Model for comparsion of other model

In [63]:

```
dataX=pd.concat([DummyManufacturer,DummyModel,DummyTransmission,DummyColor,df['odometer'],
df['year'],DummyEngineFuel,DummyEngineType,df['engineCapacity'],DummyBodyType,DummyDrivetrain,df['photos'],DummyFeature],axis=1)
dataY=df['price']
```

In [64]:

```
import sklearn.model_selection as model_selection
```

```
X_train, X_test, y_train, y_test = model_selection.train_test_split(dataX,dataY, train_size=0.8,test_size=0.2, random_state=100)
```

In [65]:

```
import numpy as np
from sklearn.linear_model import LinearRegression
model = LinearRegression().fit(X_train, y_train)
y_predict= model.predict(X_test)
import math
A=(y_predict-y_test)**2
RMSE=math.sqrt(A.sum()/y_test.size);
print(RMSE)
```

563611274.4155275

MODEL_00 linear regression with Important Variables

In [67]:

```
dataX=pd.concat([df['odometer'],df['year'],DummyManufacturer,DummyTransmission,DummyColor,
DummyBodyType,DummyDrivetrain,DummyFeature],axis=1)
dataY=df['price']
```

In [68]:

```
import sklearn.model_selection as model_selection

X_train, X_test, y_train, y_test = model_selection.train_test_split(dataX, dataY, train_size=0.8, test_size=0.2, random_state=100)

import numpy as np
from sklearn.linear_model import LinearRegression
model = LinearRegression().fit(X_train, y_train)
```

In [69]:

```
y_predict= model.predict(X_test)
import math
A=(y_predict-y_test)**2
RMSE=math.sqrt(A.sum()/y_test.size);
print(RMSE)
```

327395674.97051287

FINDING THE BEST MODEL

MODEL_01

Continuous numerical variables:

odometer => quadratic harm equation

year => quadratic harm equation

Categorical variables:

manufacturer
transmission
color
bodyType
drivetrain
feature_0
feature_1
feature_2
feature_3
feature_4
feature_5
feature_6
feature_7
feature_8

In [70]:

```
dataX=pd.concat([df['odometer'],df['odometer']**2,df['year'],df['year']**2,DummyManufacturer,DummyTransmission,DummyColor,DummyBodyType,DummyDrivetrain,DummyFeature],axis=1)
dataY=df['price']
import math

import sklearn.model_selection as model_selection

X_train, X_test, y_train, y_test = model_selection.train_test_split(dataX,dataY, train_size=0.8,test_size=0.2, random_state=100)

import numpy as np
from sklearn.linear_model import LinearRegression
model = LinearRegression().fit(X_train, y_train)

y_predicted = model.predict(X_test)
A=(y_predicted-y_test)**2
RMSE=math.sqrt(A.sum()/y_test.size)
print(RMSE)
```

2742.013157337794